

AUTONOMOUS VEHICLE MOTION PLANNING VIA PARTICLE FILTERING AND SMOOTHING

MASON LIEB

University of Kansas - Department of Mechanical Engineering
M.Sc. Project Report

M.Sc. Project Defense Committee
Huazhen Fang, Ph.D.
Sara Wilson, Ph.D.
Bozenna Pasik-Duncan, Ph.D.

Chair, Associate Professor of Mechanical Engineering
Associate Professor of Mechanical Engineering
Professor of Mathematics



Abstract

This study investigates an application of particle filtering and reweighting particle smoothing to solving a motion planning and decision-making problem for autonomous road vehicles. The method utilizes a set of driving requirements as a reference for the generated trajectory to track, where these driving requirements can be determined a priori, and no global waypoints are required. The motion planning problem is viewed as an estimation problem which is solved with particle filtering and reweighting particle smoothing. The driving requirements and therefore the reference can vary as the selected driving mode changes, which characterizes several different maneuvers a vehicle might make in a standard highway driving scenario. The decision of which driving mode to select is part of the solution framework. Collision avoidance and road boundary constraints are built into the filtering and decision-making to enhance the safety of the generated trajectories. The trajectories are applied in a receding horizon manner to facilitate applicability to dynamic environments. The specific scenario considered in this study is overtaking, with obstacle vehicles travelling at different speeds in each of the two lanes in the simulated highway. A GitHub repository with animated simulation results is included with this study. The results show the capability of the method to track reference lane centers and velocities while satisfying safety constraints with multiple driving mode changes throughout the simulation.

Contents

| | |
|---|----|
| Abstract..... | 2 |
| Contents..... | 3 |
| List of Symbols | 4 |
| 1. Introduction | 6 |
| 2. Problem Formulation | 7 |
| 3. Problem Solution | 10 |
| A. Preliminaries on Particle Filtering and Smoothing [12] | 12 |
| B. Measurement Model | 19 |
| C. Driving Mode Design and Selection | 25 |
| 4. Simulation Study | 28 |
| 5. Conclusion | 35 |
| 6. References | 37 |
| Appendix | 39 |

List of Symbols

| | |
|------------------------|--|
| p_{xk} | Global x-coordinate of the ego vehicle |
| p_{yk} | Global y-coordinate of the ego vehicle |
| ψ_k | Heading (yaw) angle of the ego vehicle [rad] |
| v_{xk} | Longitudinal velocity of the ego vehicle [m/s] |
| δ_k | Steering rate [rad/s] |
| β_k | Body slip angle (rad) |
| L | Vehicle wheelbase |
| l_r | Distance from vehicle center to vehicle rear axle |
| v_{nom} | Nominal velocity (speed limit) [m/s] |
| e_L | Lane centerline tracking error [m] |
| e_ψ | Heading angle tracking error [rad] |
| g_o | Obstacle barrier function for obstacle $o = 1, \dots, N_{obs}$ |
| g_b | Road boundary barrier function |
| Z | Driving scenario information |
| \mathbf{r} | Reference vector for measurement tracking |
| T | Planning horizon length [s] |
| Δt | Discretization time [s] |
| \mathbf{x}_k | State vector |
| n_x | Dimension of state vector |
| \mathbf{u}_k | Input vector |
| n_u | Dimension of input vector |
| $\bar{\mathbf{x}}_k$ | Virtual system state vector with augmented controls |
| $f(\cdot)$ | System dynamics $f: \mathbb{R}^{n_x+n_u} \rightarrow \mathbb{R}^{n_x+n_u}$ |
| w | Lane width [m] |
| \mathcal{M} | Set of driving modes |
| m_j | Specific driving mode instance |
| M_{max} | Maximum number of driving modes to sample in each planning phase |
| \mathbf{q}_k | Input noise |
| \mathbf{Q} | Input noise covariance matrix ($n_u \times n_u$) |
| \mathbf{Q}_0 | Process noise covariance matrix ($n_x \times n_x$) |
| $\bar{\mathbf{q}}_k$ | Virtual system process noise |
| $\bar{\mathbf{Q}}$ | Virtual system process noise covariance matrix |
| \mathbf{y}_k | Measurement vector |
| \mathbf{v}_k | Measurement noise vector |
| \mathbf{R} | Measurement noise covariance matrix |
| $h(\cdot)$ | Measurement function $h: \mathbb{R}^{n_x} \rightarrow \mathbb{R}^{n_y}$ |
| α, β | Softplus barrier function parameters for obstacle avoidance constraint |
| m, n | Softplus barrier function parameters for road boundary constraint |
| $\bar{\mathbf{x}}_k^i$ | Particle i at time k |
| N | Number of particles |
| N_{eff} | Effective number of particles |
| N_{req} | Required number of effective particles |
| \tilde{w}_k^i | Unnormalized filtering weight for particle i |
| w_k^i | Filtering weight for particle i |
| $w_{k k+1}^i$ | Smoothing weight for particle i at time k given information up to time $k + 1$ |

| | |
|------------------------------------|--|
| $\delta(\cdot)$ | Dirac delta function |
| a, b, c | Ellipse parameters |
| $d_{centers}$ | Center-to-center distance between ellipses |
| $\bar{\theta}$ | Global angle of center-to-center line between ellipses |
| \bar{r} | Specific radius of ellipse evaluated at $\bar{\theta}$ |
| Δ | Edge-to-edge distance between ellipses |
| Δ_l | Edge-to-edge distance between ego vehicle and only the obstacle in same lane |
| P_{base} | Base probability of sampling lane-keeping driving mode |
| P_{min} | Minimum probability of sampling lane-keeping driving mode |
| $\gamma, \epsilon, \zeta, \lambda$ | Weighting parameters for trajectory cost function |
| ν, η | Exponential parameters of softplus barrier functions in trajectory cost function |
| ξ, ρ | Offset terms for barrier function terms of trajectory cost function |
| G | Covariance scaling term |
| \mathbf{Q}' | Scaled virtual system process noise covariance matrix |
| \mathbf{R}' | Scaled measurement noise covariance matrix |

Notations

| | |
|-----------------|---|
| \propto | Denotes “proportional to” |
| \sim | Denotes “sampled from” |
| $(\cdot)_k$ | k – subscript indexes discretized time step |
| $(\cdot)^i$ | i – superscript indexes particles |
| $(\cdot)_{ego}$ | <i>ego</i> – subscript indicates information pertains to ego vehicle |
| $(\cdot)_{obs}$ | <i>obs</i> – subscript indicates information pertains to obstacle vehicle |
| \dot{x} | dot notation used to represent time derivatives $\dot{x} = \frac{dx}{dt}$ |

1. Introduction

There is ever-increasing interest in autonomous transportation systems in our modern society as seen in the popularity of adaptive driver assistance systems (ADAS) such as General Motors Super Cruise, Mercedes-Benz Drive Pilot, Ford BlueCruise, and Tesla Autopilot. To implement such systems, there are many different modules that must work together to receive instructions from a high-level navigation system, make decisions, perform motion planning, and execute controller inputs for low-level systems like steering, acceleration, or braking actuators. In addition to performing these actions, there are sensors that provide the necessary information to facilitate the decision making and motion planning processes. Figure 1 outlines a high-level overview of the system schematics for an autonomous vehicle.




Figure 1. Autonomous vehicle high-level system architecture [7].

The route planner sets the high-level goals of the autonomous driving system. An example of this sort of high-level route planning would be a mapping application directing a driver to take a certain exit from a highway or turn a certain direction at an intersection. The next lowest level is the decision-making system. This system determines things like which lane the vehicle should be in, which will be described by the *driving mode* in this study. The driving modes considered are to stay in the current lane, to change to the left lane, and change to the right lane. These modes capture most of the possible behaviors in a highway driving scenario apart from emergency braking, which can be handled by a separate emergency system that would also handle other evasive maneuvers. Below the decision-making level is the motion planner, which determines a trajectory that the vehicle should take based on the decision-making output and the sensor information. The generated trajectory should be collision-free and dynamically feasible while adhering to all other safety or road constraints like minimum distance to obstacles and speed limit. Vehicle control, actuator control, and final motion of the vehicle is handled by lower-level controllers, and the vehicle interacts with its dynamic environment which requires continuous sensing and mapping to determine the vehicle's position relative to the road and other vehicles or static obstacles.

The sensing and mapping in a current ADAS implementation could be based on radar and lidar sensors, stereo camera systems, wheel speed and steering angle sensors, and GPS. With the information provided from the sensing and mapping, the ego vehicle should be able to determine where it is relative to the road it is driving on, identify obstacles and their locations relative to the ego vehicle, as well as identifying the road constraints such as speed limit.

Motion planning for autonomous road vehicles has several different classes of solutions, including graph search-based planners, interpolating curve planners, numerical optimization, and sampling-based planners [2]. Example implementations of each class of planning methods can be found in [2], as well as descriptions of their advantages and disadvantages. For this study, there are inspirations drawn from the numerical optimization and sampling-based classes of motion planning solutions.

An example of an optimization-based solution method is model predictive control (MPC), where over some finite horizon, a trajectory is optimized over the inputs to the system. These inputs are then applied

in a receding horizon manner, where only the first generated input is applied to the true system while the optimization repeats over the next horizon. MPC has been applied to robotic motion planning problems for legged robots [4] and autonomous vehicles [13]. A challenge facing the MPC method for autonomous vehicles is nonconvexity, which may cause the optimization to fall into local minima and fail to find an optimal motion plan.

From a sampling-based perspective, one can sample configurations or inputs. Rapidly exploring random trees were originally proposed as a method of computing collision-free kinodynamic trajectories for high degree-of-freedom problems [11]. RRTs were applied to autonomous driving motion planning in [14]. Since their introduction, they have been expanded into the probabilistically complete RRT* [10] where, as the number of samples approaches infinity, a feasible trajectory (if one exists) will be found with probability one. These methods sample the configuration space and will generate satisfactory paths, but for an autonomous vehicle, the steering inputs must then be calculated via inverse steering dynamics. By constructing a virtual system that augments the system inputs with the states, sampling the augmented state space also samples the control inputs. Sampling the control inputs is a smaller sample space and avoids the calculation of inverse steering laws. And, by propagating the samples through the system's dynamic model, the resulting trajectories are guaranteed to be dynamically feasible. This study does not seek to find an optimal trajectory, as in normal driving there are many trajectories that are suitable and satisfy the driving requirements and there is not a general requirement to minimize the amount of time a maneuver takes. Thus, the trajectories that are developed with this method are not necessarily optimal but are suitable and safe.

To accomplish this, the samples drawn are used as particles in the particle filtering framework. The generated trajectories from filtering are then smoothed via a reweighting particle smoother before the inputs are extracted and applied in a receding horizon manner. Using particle filtering to perform MPC was applied in [3] [5]. The inspiration for this study is the particle filter motion planning and decision-making method presented in [6] [7].

The remainder of the report is organized as follows. First, the high-level motion planning problem is outlined with the driving scenario considered, the vehicle model, measurement model, and the driving modes for decision-making. Then, the particle filtering and smoothing solution methodology is described. Next, a detailed description for each measurement and the driving mode selection scheme is given. Finally, simulation results are displayed before concluding remarks.

2. Problem Formulation

Driving Scenario

This study considers the problem of generating control inputs that drive a vehicle such that its trajectories satisfy a given reference that includes road requirements such as speed limit and safety constraints like road boundary satisfaction and obstacle avoidance. The driving scenario is on a straight two-lane highway, with dynamic obstacles. Under normal highway driving conditions, acceleration and steering inputs are relatively small, which means that the executed maneuvers are well within the vehicle's performance envelope. The motion planning module of an autonomous vehicle would feed control inputs to lower-level controllers for the vehicle actuators, which in a practical application would handle the errors between the kinematic model used and the true vehicle dynamics. It is assumed that the ego vehicle can measure its own position relative to the road boundaries and lane centerlines, as well

as the positions and velocities of the obstacle vehicles. Other available information includes the speed limit and the ego vehicle's heading angle relative to the road.

Vehicle Model

This study considers the kinematic bicycle (or single-track) model as it provides satisfactory accuracy for motion planning purposes. The equations of motion for this model are

$$\dot{x} = \begin{bmatrix} \dot{p}_x \\ \dot{p}_y \\ \dot{\psi} \\ \dot{v}_x \\ \dot{\delta} \end{bmatrix} = \begin{bmatrix} \frac{v_x \cos(\psi + \beta)}{\cos(\beta)} \\ \frac{v_x \sin(\psi + \beta)}{\cos(\beta)} \\ \frac{v_x \tan(\delta)}{L} \\ u_1 \\ u_2 \end{bmatrix} \quad (1)$$

$$\beta = \arctan\left(\frac{l_r \tan(\delta)}{L}\right),$$

where β is the kinematic body slip angle, L is the wheelbase of the vehicle and l_r is the distance from the center of the vehicle to the rear axle [7]. Dot notation is used to indicate the time derivative. The vehicle state variables are the vehicles x- and y-positions (p_x, p_y), yaw (heading) angle (ψ), longitudinal velocity (v_x), and steering angle (δ). The state variables are illustrated below in Figure 2. The inputs to the system are acceleration and steering rate, which guarantees smooth trajectories when the inputs are applied to the actual system.




Figure 2. Illustration of state variables for kinematic bicycle model.

The model is discretized as below to be compatible with the sequential estimation procedure of the particle filter which is described in detail in Section 3A. The subscript k indicates a discrete time index.

$$x_{k+1} = x_k + \dot{x}_k \Delta t$$

$$\begin{bmatrix} p_x \\ p_y \\ \psi \\ v_x \\ \delta \end{bmatrix}_{k+1} = \begin{bmatrix} p_{x_k} + \frac{v_{x_k} \cos(\psi_k + \beta_k)}{\cos(\beta_k)} \Delta t \\ p_{y_k} + \frac{v_{x_k} \sin(\psi_k + \beta_k)}{\cos(\beta_k)} \Delta t \\ \psi_k + \frac{v_{x_k} \tan(\delta_k)}{L} \Delta t \\ v_{x_k} + u_{1_k} \Delta t \\ \delta_k + u_{2_k} \Delta t \end{bmatrix} = f(x_k, u_k) \quad (2)$$

Reference Tracking

The reference that the vehicle should track needs to satisfy several requirements for generated trajectories to be safe and efficient. The first set of requirements can be thought of as the road requirements, which are determined by the road itself. These requirements include tracking the road's speed limit, denoted as v_{nom} or nominal velocity, track the centerline of the vehicle's target lane, do not leave the road, and maintain no deviation in heading angle from the road angle. The second set of requirements concerns the obstacles present in the scenario, namely the other vehicles driving on the road. The only requirement concerning the obstacle vehicles is that the ego vehicle maintains a minimum distance from them that is a function of the ego vehicle's velocity. This means that the faster the ego vehicle is travelling, the larger the collision avoidance zone will be. This is illustrated below in Figure 3.




Figure 3. Obstacle region with ego vehicle speed of 20 m/s (top). Obstacle region with ego vehicle speed of 30 m/s (bottom)

This reference is assumed fixed during each planning horizon, so only a single realization of the reference vector is required, denoted as \mathbf{r} . The reference vector is described below in (3).

$$\mathbf{r} = [v_{nom} \quad e_L = 0 \quad g_o = 0 \quad g_b = 0 \quad e_\psi = 0]^T \quad (3)$$

The reference includes four terms regarding the road itself: nominal velocity or speed limit v_{nom} , lateral lane centerline tracking error e_L , road boundary satisfaction g_b , and heading angle error e_ψ . Obstacle avoidance is enabled by generating N_{obs} obstacle reference terms, which are described by g_o where $o =$

$1, \dots, N_{obs}$. For a single obstacle, \mathbf{r} will have 5 elements. For two obstacles, \mathbf{r} will have 6 elements, and so on.

Driving Modes

Distinct driving modes are required to determine the reference \mathbf{r} . The driving modes will tell the vehicle what velocity and lane centerline to track. Driving modes can be selected in different ways. One method is to consider a feasible set of driving modes \mathcal{M} and generate a single trajectory according to each driving mode. Another method is by generating samples from a transition probability $m_j \sim p(m_j|Z)$ of switching to mode j conditioned on the currently available road and obstacle information Z . The notation " \sim " reads as "is sampled from". With either the deterministic or probabilistic method, the "best" trajectory out of those generated will be determined by a cost function and executed in a receding horizon manner.

The benefit of deterministically testing each viable driving mode is that each mode is guaranteed to be checked. This contrasts with the transition probability sampling, where it is possible that a feasible mode may not be checked during trajectory generation unless there is a large enough amount of samples, which comes with significant computational cost. However, if the transition probabilities are designed well and appropriately conditioned on currently available information, then that method may lead to a more optimal trajectory. This is because the mode that is expected to produce high quality trajectories will get sampled more, and more possible trajectories corresponding to that mode will be generated. The probabilistic sampling method is used to obtain the results of this study.

The considered driving modes are lane-keeping, change to the left lane, and change to the right lane. These are abbreviated and described in the set $\mathcal{M} = \{LK, CLL, CLR\}$. The nominal velocity for all driving modes is the speed limit. For lane-keeping, the target lane is the current lane. For the lane-changing modes, the target lane is one lane left or right of the current lane. In a practical sense, lane changing is shifting the reference lane by w meters, where w is the width of a lane. The other elements of the reference signal (obstacle avoidance, road boundary constraint, and heading angle constraint) are all left unchanged for all driving modes, as those elements always need to be consistently satisfied. Each driving mode can have varied parameters that affect the generation of the corresponding trajectory, and these parameters are outlined in Section 3C. Section 4 includes discussion of the specific parameter values used to obtain the results of this study.

With the system and measurement model described here, the problem being addressed is how to generate a viable sequence of control inputs $u_{k:k+T}$ over the planning horizon that drives the system towards the reference signal. In the next section, the specific solution to this problem used in this study is given.

3. Problem Solution

Overview

Particle filtering and smoothing are used to address the motion planning problem by first estimating a trajectory over the planning horizon and then applying the generated inputs from said trajectory in a receding horizon fashion. This method is like MPC, where an optimization over the inputs is performed over the planning horizon under a set of constraints, and then those inputs are applied in a receding horizon manner. Equivalence between the solution to a nonlinear model-predictive control problem via constrained optimization and via estimation is described further in [5]. The constraints in this problem are

present in the design of the reference signal and the corresponding covariance matrix for the measurement noise. This implementation results in only soft constraints being present where there are only probabilistic guarantees of constraint satisfaction. These constraints are made harder via the implementation of barrier functions in the measurement function $h(\cdot)$.

The goal is to estimate a posterior distribution of states and inputs given the reference \mathbf{r} .

$$p(\mathbf{x}_{k:k+T}, \mathbf{u}_{k:k+T} | \mathbf{r}) \quad (4)$$

The reference \mathbf{r} does not carry a subscript because it is constant over each planning horizon. The distribution in (4) describes a joint posterior over the states and inputs over a horizon of length T .

To convert the model into one that lends itself to the sequential estimation of the particle filter, the controls are modeled as random noise acting on the system. It is desired that the noise be additive, as this simplifies the process modeling. The system is rewritten as follows, starting from the previously described system dynamics in (2).

$$\mathbf{x}_{k+1} = f(\mathbf{x}_k, \mathbf{u}_k) \quad (5)$$

The inputs are replaced with random process noise \mathbf{q}_k . Then the system is rewritten in terms of the augmented state vector $\bar{\mathbf{x}}_k$.

$$\mathbf{u}_k = \mathbf{q}_k \quad (6)$$

$$\bar{\mathbf{x}}_{k+1} = f(\bar{\mathbf{x}}_k) + \bar{\mathbf{q}}_k \quad (7)$$

$$\bar{\mathbf{x}}_k = \begin{bmatrix} \mathbf{x}_k \\ \mathbf{u}_k \end{bmatrix} = \begin{bmatrix} p_x \\ p_y \\ \psi \\ v_x \\ \delta \\ \dot{v}_x \\ \dot{\delta} \end{bmatrix}_k, \bar{\mathbf{q}}_k = \begin{bmatrix} \mathbf{0}_{n_x \times 1} \\ \mathbf{q}_k \end{bmatrix} \quad (8)$$

In this *virtual system*, the mean and covariance of the process noise \mathbf{q}_k can be designed to constrain the generated inputs. The noise is modeled as white Gaussian, hence has zero mean and a $n_u \times n_u$ covariance matrix \mathbf{Q} , where n_u is the dimension of the control input space (in this study, $n_u = 2$).

$$\mathbf{q}_k \sim \mathcal{N}(\mathbf{0}, \mathbf{Q}) \quad (9)$$

Thus, the virtual system also follows a Gaussian distribution with mean $\bar{\mathbf{x}}_k$ and covariance $\bar{\mathbf{Q}}$, where $\bar{\mathbf{Q}}$ is a block diagonal matrix constructed as follows.

$$\bar{\mathbf{Q}} = \begin{bmatrix} \mathbf{Q}_0 & \mathbf{0} \\ \mathbf{0} & \mathbf{Q} \end{bmatrix} \quad (10)$$

Here, the lower right submatrix is the tuned covariance for the inputs and the upper $n_x \times n_x$ submatrix \mathbf{Q}_0 is diagonal with each element set to be approximately zero. This is done to facilitate numerical stability for the reweighting smoother, discussed later in this section. It also implies that the system has a high level of confidence in the evolution of the state variables. The covariance for the inputs can vary according to the driving mode, however the process noise for the state variables of the original system (\mathbf{Q}_0) is kept constant.

The measurements \mathbf{y}_k are designed to evaluate how well the current vehicle augmented state satisfies the driving requirements in the reference vector \mathbf{r} .

$$\mathbf{y}_k = \mathbf{h}(\mathbf{x}_k) + \mathbf{v}_k \quad (11)$$

The measurement noise \mathbf{v}_k is taken to be white Gaussian with covariance \mathbf{R} , which is designed to put more emphasis on certain driving requirements for safety purposes.

$$\mathbf{v}_k \sim \mathcal{N}(\mathbf{0}, \mathbf{R}) \quad (12)$$

For example, the covariance term for velocity tracking could be set to be relatively large, which allows for a decent amount of deviation from the reference velocity. This is balanced by setting tight constraints on obstacle avoidance or lane centerline tracking. The measurement function $\mathbf{h}(\cdot)$ is summarized below.

$$\mathbf{h}(\mathbf{x}_k) = \begin{bmatrix} \mathbf{x}_k(4) \\ e_L \\ g_o(\mathbf{x}_k) \\ g_b(\mathbf{x}_k) \\ e_\psi \end{bmatrix} = \begin{bmatrix} \text{longitudinal velocity } v_{x_k} \\ \text{lateral lane center tracking error} \\ \text{obstacle barrier function(s)} \\ \text{road boundary barrier function} \\ \text{heading angle tracking error} \end{bmatrix} \quad (13)$$

The control inputs are not considered in the measurement function, so the original state vector \mathbf{x}_k can be used as the input argument. The design of the functions that constitute $\mathbf{h}(\cdot)$ is described in detail in Section 3B. To give context to how these measurements are used to evaluate how well a state satisfies the reference, the estimation methodology is first described in detail.

A. Preliminaries on Particle Filtering and Smoothing [12]

Forward Filtering

Particle filtering is an estimation method that can achieve arbitrarily good estimates and can handle nonlinear and potentially multimodal non-Gaussian signals. Further detailed derivation of the particle filter used in this study is provided in the Appendix, but the main filtering equations for sequential importance sampling/resampling are covered here.

Particle filtering traces to the Monte Carlo method, a general method of finding expectations over the posterior distribution.

$$\mathbb{E}[g(\mathbf{x})|\mathbf{y}_{1:T}] = \int g(\mathbf{x})p(\mathbf{x}|\mathbf{y}_{1:T})d\mathbf{x} \quad (14)$$

Here, $g(\mathbf{x})$ is an arbitrary function of the random variable \mathbf{x} which represents the latent state vector of the system. This integral generally does not have a closed-form solution, so Monte Carlo sampling-based methods are a common choice for solving this type of problem. For the integral in (14), if one could draw samples directly from the posterior distribution, then the evaluation of the integral could be approximated as

$$\mathbb{E}[g(\mathbf{x})|\mathbf{y}_{1:T}] = \int g(\mathbf{x})p(\mathbf{x}|\mathbf{y}_{1:T})d\mathbf{x} \approx \frac{1}{N} \sum_{i=1}^N g(\mathbf{x}^i), \quad (15)$$

where the convergence of this approximation is guaranteed by the central limit theorem [12]. Unfortunately, the true posterior distribution may be difficult to draw samples from, which motivates the use of importance sampling. Importance sampling is a method of drawing samples from an importance (or proposal) distribution. Then, by applying weights to those samples determined by how well that

sample matches the target distribution (the true posterior), the target distribution can be approximated as

$$\begin{aligned} \int g(\mathbf{x})p(\mathbf{x}|\mathbf{y}_{1:T})d\mathbf{x} &= \int \left[g(\mathbf{x}) \frac{p(\mathbf{x}|\mathbf{y}_{1:T})}{\pi(\mathbf{x}|\mathbf{y}_{1:T})} \right] \pi(\mathbf{x}|\mathbf{y}_{1:T})d\mathbf{x} \\ &\approx \frac{1}{N} \sum_{i=1}^N \frac{p(\mathbf{x}|\mathbf{y}_{1:T})}{\pi(\mathbf{x}|\mathbf{y}_{1:T})} g(\mathbf{x}^i) = \sum_{i=1}^N \tilde{w}^i g(\mathbf{x}^i), \end{aligned} \quad (16)$$

where $\mathbf{x}^i \sim \pi(\mathbf{x}|\mathbf{y}_{1:T})$ and the superscript $(\cdot)^i$ denotes a particle index. Here, the weights are defined as

$$\tilde{w}^i = \frac{1}{N} \frac{p(\mathbf{x}^i|\mathbf{y}_{1:T})}{\pi(\mathbf{x}^i|\mathbf{y}_{1:T})}. \quad (17)$$

An illustrative example of importance sampling is shown below with a simple case of a standard normal distribution as the target and a uniform distribution over the range of [-5,5] acting as the proposal.




Figure 4. Importance Sampling Example. The circles represent the approximation formed by the normalized weighted samples.

Problems can arise with this method due to difficulty in evaluating the full posterior at each sample $p(\mathbf{x}^i|\mathbf{y}_{1:T})$. Bayes' rule can be applied to the posterior distribution to provide a different method of arriving at an expression for normalized weights while avoiding this evaluation. First, Bayes' rule is applied to the posterior distribution for the samples.

$$p(\mathbf{x}^i|\mathbf{y}_{1:T}) = \frac{p(\mathbf{y}_{1:T}|\mathbf{x}^i)p(\mathbf{x}^i)}{p(\mathbf{y}_{1:T})} = \frac{p(\mathbf{y}_{1:T}|\mathbf{x}^i)p(\mathbf{x}^i)}{\int p(\mathbf{y}_{1:T}|\mathbf{x})p(\mathbf{x})d\mathbf{x}} \quad (18)$$

The terms in the numerator can generally be evaluated, but another importance sampling approximation is required to compute the marginal likelihood term in the denominator.

$$\begin{aligned}
\mathbb{E}[g(\mathbf{x})|\mathbf{y}_{1:T}] &= \int g(\mathbf{x})p(\mathbf{x}|\mathbf{y}_{1:T})d\mathbf{x} = \frac{\int g(\mathbf{x})p(\mathbf{y}_{1:T}|\mathbf{x})p(\mathbf{x})d\mathbf{x}}{\int p(\mathbf{y}_{1:T}|\mathbf{x})p(\mathbf{x})d\mathbf{x}} \\
&= \frac{\int \left[\frac{p(\mathbf{y}_{1:T}|\mathbf{x})p(\mathbf{x})}{\pi(\mathbf{x}|\mathbf{y}_{1:T})} \right] g(\mathbf{x})\pi(\mathbf{x}|\mathbf{y}_{1:T})d\mathbf{x}}{\int \left[\frac{p(\mathbf{y}_{1:T}|\mathbf{x})p(\mathbf{x})}{\pi(\mathbf{x}|\mathbf{y}_{1:T})} \right] \pi(\mathbf{x}|\mathbf{y}_{1:T})d\mathbf{x}} \\
&\approx \frac{\frac{1}{N} \sum_{i=1}^N \frac{p(\mathbf{y}_{1:T}|\mathbf{x}^i)p(\mathbf{x}^i)}{\pi(\mathbf{x}^i|\mathbf{y}_{1:T})} g(\mathbf{x}^i)}{\frac{1}{N} \sum_{j=1}^N \frac{p(\mathbf{y}_{1:T}|\mathbf{x}^j)p(\mathbf{x}^j)}{\pi(\mathbf{x}^j|\mathbf{y}_{1:T})}} \\
&= \sum_{i=1}^N \left[\frac{\frac{p(\mathbf{y}_{1:T}|\mathbf{x}^i)p(\mathbf{x}^i)}{\pi(\mathbf{x}^i|\mathbf{y}_{1:T})}}{\sum_{j=1}^N \frac{p(\mathbf{y}_{1:T}|\mathbf{x}^j)p(\mathbf{x}^j)}{\pi(\mathbf{x}^j|\mathbf{y}_{1:T})}} \right] g(\mathbf{x}^i) = \sum_{i=1}^N w^i g(\mathbf{x}^i)
\end{aligned} \tag{19}$$

With this process, the final expression for the sample weights is normalized and only requires evaluation of the measurement likelihood and the prior for each sample. The posterior distribution can then be approximated by the samples and their weights by (20) below:

$$p(\mathbf{x}|\mathbf{y}_{1:T}) \approx \sum_{i=1}^N w^i \delta(\mathbf{x} - \mathbf{x}^i) \tag{20}$$

Here, $\delta(\cdot)$ is the Dirac delta function. Up to this point, the estimated density is the full posterior. If measurements are obtained at each time step, which is the case in this study, then the filtering density needs to be updated sequentially as new information becomes available. This requires extending the above method to sequential importance sampling.

Sequential importance sampling in general considers Markov chain systems that have dynamics of the following form:

$$\begin{aligned}
\mathbf{x}_k &\sim p(\mathbf{x}_k|\mathbf{x}_{k-1}) \\
\mathbf{y}_k &\sim p(\mathbf{y}_k|\mathbf{x}_k)
\end{aligned} \tag{21}$$

Instead of estimating the full posterior $p(\mathbf{x}|\mathbf{y}_{1:T})$ at each time step, the filtering density $p(\mathbf{x}_k|\mathbf{y}_{1:k})$ can be estimated and used to recursively update the full posterior as follows.

$$\begin{aligned}
p(\mathbf{x}_{0:k}|\mathbf{y}_{1:k}) &\propto p(\mathbf{y}_k|\mathbf{x}_{0:k}, \mathbf{y}_{1:k-1})p(\mathbf{x}_{0:k}|\mathbf{y}_{1:k-1}) \\
&= p(\mathbf{y}_k|\mathbf{x}_k)p(\mathbf{x}_k|\mathbf{x}_{0:k-1}, \mathbf{y}_{1:k-1})p(\mathbf{x}_{0:k-1}|\mathbf{y}_{1:k-1}) \\
&= p(\mathbf{y}_k|\mathbf{x}_k)p(\mathbf{x}_k|\mathbf{x}_{k-1})p(\mathbf{x}_{0:k-1}|\mathbf{y}_{1:k-1})
\end{aligned} \tag{22}$$

The notation “ \propto ” reads as “is proportional to”. With this recursive representation of the posterior distribution, a proposal distribution $\pi(\cdot)$ needs to be developed to develop the recursive weight update (25). Carrying the idea of defining weights as a ratio of posterior to proposal, the weight update is expanded using (22) as

$$\begin{aligned}
w_k^i &\propto \frac{p(\mathbf{y}_k|\mathbf{x}_k^i)p(\mathbf{x}_k^i|\mathbf{x}_{k-1}^i)p(\mathbf{x}_{0:k-1}^i|\mathbf{y}_{1:k-1})}{\pi(\mathbf{x}_{0:k}^i|\mathbf{y}_{1:k})} \\
&= \frac{p(\mathbf{y}_k|\mathbf{x}_k^i)p(\mathbf{x}_k^i|\mathbf{x}_{k-1}^i)}{\pi(\mathbf{x}_k^i|\mathbf{x}_{0:k-1}^i,\mathbf{y}_{1:k})} \times \frac{p(\mathbf{x}_{0:k-1}^i|\mathbf{y}_{1:k-1})}{\pi(\mathbf{x}_{0:k-1}^i|\mathbf{y}_{1:k-1})}.
\end{aligned} \tag{23}$$

To decompose the proposal distribution as in (23), the following property (24) must be enforced. Further details on the derivation of (23) are presented in the Appendix.

$$\pi(\mathbf{x}_{0:k-1}^i|\mathbf{y}_{1:k}) = \pi(\mathbf{x}_{0:k-1}^i|\mathbf{y}_{1:k-1}) \tag{24}$$

The unnormalized weights are then updated with (25) and normalized with (26).

$$\tilde{w}_k^i = \frac{p(\mathbf{y}_k|\mathbf{x}_k^i)p(\mathbf{x}_k^i|\mathbf{x}_{k-1}^i)}{\pi(\mathbf{x}_k^i|\mathbf{x}_{0:k-1}^i,\mathbf{y}_{1:k})} \times w_{k-1}^i \tag{25}$$

$$w_k^i = \frac{\tilde{w}_k^i}{\sum_i \tilde{w}_k^i} \tag{26}$$

Now that the weights can be updated sequentially once new information y_k arrives, the filtering distribution can also be updated sequentially and is represented as shown in (27) below:

$$p(\mathbf{x}_k|\mathbf{y}_{1:k}) \approx \sum_{i=1}^N w_k^i \delta(\mathbf{x}_k - \mathbf{x}_k^i) \tag{27}$$

For this study, a variation of the particle filter called the *bootstrap particle filter* is used, where the proposal distribution is chosen to be the system dynamics.

$$\pi(\mathbf{x}_k|\mathbf{x}_{0:k-1},\mathbf{y}_{1:k}) = p(\mathbf{x}_k|\mathbf{x}_{k-1}) \tag{28}$$

This simplifies the weight update to be proportional to just the measurement likelihood.

$$w_k^i \propto w_{k-1}^i p(\mathbf{y}_k|\mathbf{x}_k^i) \tag{29}$$

Adapting this development of the bootstrap particle filter to the notation of the problem at hand, the measurements are taken to be the reference that the vehicle needs to track \mathbf{r} , and the states will be the augmented state and input vector from the virtual system. The states are mapped to the reference space by the measurement function. The weights are then updated with (30), normalized with (26) and applied to the particle filter representation of the posterior distribution in (31).

$$\tilde{w}_k^i = w_{k-1}^i p(\mathbf{r}|\bar{\mathbf{x}}_k^i) \tag{30}$$

$$p(\bar{\mathbf{x}}_k|\mathbf{r}) \approx \sum_{i=1}^N w_k^i \delta(\bar{\mathbf{x}}_k - \bar{\mathbf{x}}_k^i) \tag{31}$$

One way to think about this implementation is that a state that satisfies the reference signal well will have a higher measurement likelihood and therefore a larger weight than a state that poorly satisfies the reference.

Resampling

As the state evolves over time, the estimation quality can degrade due to a phenomenon called particle degeneracy. This happens when only a few of the overall particles have appreciable weights, which makes the estimate of the posterior distribution poor. Resampling the particles is a method to combat the particle degeneracy problem by replacing bad or low-weight particles. There are many different methods of resampling [1], and in this study traditional multinomial resampling is used. Multinomial resampling is illustrated in Figure 5 where a uniform distribution is sampled for an index i , and then the particle that is resampled is indexed as j . The cumulative distribution function in Figure 5 is constructed with the weighted set of particles. The particles with larger weights are more likely to get resampled, which will replace the particles with low weights.




Figure 5. Multinomial resampling illustration [9]

The resampling procedure will increase the computational effort required by the algorithm, so it is not performed at every step. Whether to do resampling or not is based on the number of effective particles, which is calculated as

$$N_{eff} = \frac{1}{\sum_{i=1}^N (w_k^i)^2}. \quad (32)$$

A threshold value of effective particles required N_{req} is set, and if the number calculated with (32) is less than this value then resampling is performed. The implementation of the bootstrap particle filter is described in Algorithm 1 below.

Algorithm 1 – Particle Filtering

- Given the constant reference \mathbf{r} for the planning horizon of length T , and current road information Z
- Sample the initial set of particles from the prior distribution $\mathbf{x}_0^i \sim p(\mathbf{x})$
- For $k = 1$ to T
 - Predict obstacle regions, update road information Z
 - For $i = 1$ to N
 - Generate $\bar{\mathbf{x}}_k^i$ with (21)
 - Update weight with (30)
 - End for
 - Normalize weights with (26)
 - Calculate N_{eff} with (32)
 - If $N_{eff} < N_{req}$
 - Resample particles with replacement
 - Set weights to be uniform
 - End if
- End for
- Pass generated sequence of weighted particles to reweighting smoother

Reweighting Particle Smoother

At the end of the planning horizon, when all measurements have been obtained, the quality of the estimation can be improved by smoothing due to the additional information available. The goal of Bayesian smoothing is to compute the marginal posterior distribution of \mathbf{x}_k given the measurements in a set window (horizon) where $T > k$.

$$p(\mathbf{x}_k | \mathbf{y}_{1:T}) \quad (33)$$

The reweighting particle smoother computes new weights for the particles to get an approximation of this smoothing distribution. As with the forward filtering equations, the main results are shown here, and the derivations are included in the Appendix.

Given the result of the forward filtering procedure, there are approximations of the marginal distributions $p(\mathbf{x}_k | \mathbf{y}_{1:k})$ and the complete measurement set $\mathbf{y}_{1:T}$ available. The smoothing weights are initialized as the terminal filtering weights

$$w_{T|T}^i = w_T^i \quad \forall i = 1, \dots, N \quad (34)$$

where the notation $w_{k|T}^i$ indicates the smoothing weight for particle i at time k given information up to time $T > k$. After the smoothing weights are initialized, they can be updated recursively backwards through the horizon.

$$w_{k|T}^i = \sum_j w_{k+1|T}^j \frac{w_k^i p(\mathbf{x}_{k+1}^i | \mathbf{x}_k^i)}{[\sum_l w_k^l p(\mathbf{x}_{k+1}^l | \mathbf{x}_k^l)]} \quad (35)$$

With these weights, the marginal distribution at each time step can be approximated as

$$p(\mathbf{x}_k | \mathbf{y}_{1:T}) \approx \sum_i w_{k|T}^i \delta(\mathbf{x}_k - \mathbf{x}_k^i). \quad (36)$$

Again, adapting these equations to the notation of the system in this study, the smoothing density is for the augmented state vector $\bar{\mathbf{x}}_k$ given the reference \mathbf{r} with the newly generated smoothing weights.

$$p(\bar{\mathbf{x}}_k | \mathbf{r}) \approx \sum_{i=1}^N w_{k|T}^i \delta(\bar{\mathbf{x}}_k - \bar{\mathbf{x}}_k^i) \quad (37)$$

With these equations, the reweighting smoother procedure is defined in the algorithm below.

Algorithm 2 – Reweighting Particle Smoothing

- Given the particles and weights from the particle filtering procedure and the complete measurement set over the horizon
- Initialize the smoothing weights with (34)
- For $k = T - 1$ to 0
 - For $i = 1$ to N
 - Calculate smoothing weights using (35)
 - End for
 - Generate marginal smoothing distribution with (37)
- End for

Once the smoothing distribution is generated, the smoothed vehicle state estimate can then be extracted and then combined into a trajectory over the planning horizon.

$$\bar{\mathbf{x}}_k = \sum_{i=1}^N w_{k|T}^i \bar{\mathbf{x}}_k^i \quad (38)$$

$$\bar{\mathbf{x}}_{0:T} = [\bar{\mathbf{x}}_0, \bar{\mathbf{x}}_1, \dots, \bar{\mathbf{x}}_{T-1}, \bar{\mathbf{x}}_T]$$

The effects of this smoothing on the trajectory generation can be seen below in Figure 6. The resulting trajectory begins the maneuver sooner and is overall smoother, as expected.




Figure 6. Filtering and smoothing trajectory lateral positions compared during a lane change maneuver. The trajectory is a single one-second planning horizon.

To drive the simulated vehicle along the generated smoothing trajectory, the inputs are extracted from the augmented states and applied in a receding horizon fashion where only the first input is used, and the rest are discarded. As the virtual system vector includes the inputs as the 6th and 7th elements, these inputs are obtained by taking the 6th and 7th elements from $\bar{\mathbf{x}}_0$ as found in (38) above. With the estimation method fully developed, the specific measurements that will be used are described in detail in the following section.

B. Measurement Model

Nominal Velocity Tracking

The velocity tracking term is direct state feedback of the vehicle's longitudinal velocity. Recall that this is the 4th element of the state vector. As the reference value for this measurement term is the nominal velocity, the two values can be directly compared.

Lane Tracking

In the considered straight road scenario, the road axis can be aligned with the global x-direction which allows for simple implementation of quantifying the lane tracking error. As shown in Figure 7 below, the lane centerlines can be described by a single global y-coordinate, to which the vehicle's global y-position p_Y can be directly compared.




Figure 7. Lane centerline illustration for straight road case.

Obstacle Avoidance

The next set of measurements are for obstacle avoidance, where there are N_{obs} terms. To model the obstacles in a parametric manner that is computationally efficient, an ellipse is generated around the obstacle and the ego vehicle, each with known functional forms. The obstacle ellipse should be designed such that the obstacle itself is completely enclosed by the region. The region is extended within the obstacle lane to provide a safety buffer for passing maneuvers to incentivize the faster-moving ego vehicle to change lanes sooner when approaching the obstacle from behind. The general polar form of an ellipse is given as

$$r(\theta) = \frac{ab}{\sqrt{(b\cos\theta)^2 + (a\sin\theta)^2}}, \quad (39)$$

where the parameters of the ellipse (a and b) are the major and minor radii, respectively. The length of the ellipse is parametrized by the ego vehicle velocity, so the obstacle region elongates as the velocity increases as shown in Figure 3. To perform this parametrization, the foci-to-center distance c is modeled as a function of EV velocity in (40) and the major radius is then calculated in (41).

$$c = 0.8v_x \quad (40)$$

$$a = \sqrt{b^2 + c^2} \quad (41)$$

The leading coefficient in (40) is designed to give the distance between the obstacle and the trailing edge of the ellipse a length of approximately v_x meters. The chosen value accomplishes this because for highway speeds, the ellipse is elongated to the point where $c \approx a$, and setting the coefficient as such results in an ellipse approximately $1.5v_x$ in overall length. The minor radius b is simply the lane half-width to make the ellipse occupy the entire width of the obstacle lane. Figure 8 displays the ellipse's parameters and its relation to the ego vehicle velocity generated by (40).




Figure 8. Elliptical obstacle region with parameters b and c shown.

As $a \approx c$, a is not shown separately in Figure 8 ($c = 24$, $a = 24.0674$). The adjustment of the ellipse center via (42) results in the distance from the obstacle to the trailing edge being approximately $v_x \frac{m}{s} * 1s$ long.

$$center = x_{obs} + \frac{v_x}{2} - c \quad (42)$$

Here, x_{obs} denotes the center of the obstacle in the global longitudinal direction, or in this case, the global x-coordinate of the obstacle center. Note that if $c = 0.75v_x$, this results in a distance between the obstacle center and the trailing edge of the ellipse that is exactly equal to the ego velocity. However, the distance that matters is to the rear of the vehicle. This motivates increasing the factor to 0.8, which generates an ellipse with the desired properties that is shown in Figure 8.

An ellipse is also used for modeling the ego vehicle's collision region, where the major and minor radii are specified by the vehicle's dimensions. For the straight road scenario considered in this study, the obstacle vehicles maintain a heading angle of zero, so the ellipses that describe their obstacle regions do not need any rotation. However, the ego vehicle's heading angle ψ will change, so the ellipse that bounds it must be rotated accordingly. In the following equations, the subscript *ego* indicates the ego vehicle, and the subscript *obs* indicates the obstacle vehicle. The bounding ellipse for an ego vehicle state is described in (43) below.

$$r_{ego}(\theta) = \frac{a_{ego}b_{ego}}{\sqrt{(b_{ego}\cos(\theta + \psi))^2 + (a_{ego}\sin(\theta + \psi))^2}} \quad (43)$$

In the particle filter motion planning implementation, each particle has an ellipse that is used in the measurement function to evaluate the particle's weight. To calculate the distance between the ellipse edges, the center-to-center distance and angle between the ellipse centers $\bar{\theta}$ is calculated.

$$d_{centers} = \sqrt{(y_{ego} - y_{obs})^2 + (x_{ego} - x_{obs})^2} \quad (44)$$

$$\bar{\theta} = \tan^{-1}\left(\frac{y_{ego} - y_{obs}}{x_{ego} - x_{obs}}\right) \quad (45)$$

Then, the *specific radius* for each ellipse is evaluated.

$$\bar{r}_{ego} = \frac{a_{ego}b_{ego}}{\sqrt{(b_{ego}\cos(\bar{\theta} + \psi))^2 + (a_{ego}\sin(\bar{\theta} + \psi))^2}} \quad (46)$$

$$\bar{r}_{obs} = \frac{a_{obs}b_{obs}}{\sqrt{(b_{obs}\cos(\bar{\theta}))^2 + (a_{obs}\sin(\bar{\theta}))^2}} \quad (47)$$

These specific radii are illustrated below in Figure 9.




Figure 9. Passing maneuver stages with scenario ellipses. The circles lying on each ellipse mark the specific radii. The small, dashed line between the ellipse centers is the center-to-center distance. The heavy dashed line is the edge-to-edge distance Δ .

Calculating the edge-to-edge distance between the ellipses is straightforward once the specific radii are calculated.

$$\Delta = d_{centers} - \bar{r}_{ego} - \bar{r}_{obs} \quad (48)$$

It can be seen in Figure 9 that this method does not return the actual minimum distance between the two ellipse edges, which is especially obvious in the third stage. Because the obstacle region is conservative, the approximation made with this method of finding obstacle distance is acceptable.

Once Δ is calculated, there remains the evaluation of the measurement function term g_o . To enforce strict adherence to obstacle avoidance requirements, a barrier function is used that will give particles in violation of the obstacle region very low weights. The softplus barrier function is chosen.

$$g(\Delta) = \frac{1}{\alpha} \log(1 + e^{\beta\Delta}) \quad (49)$$

This function has two parameters α and β which, if chosen properly, cause the function output to be zero for positive arguments and grow rapidly for negative arguments. This can be thought of as an inequality constraint, where the requirement is for $g(x) \geq 0$. To get this behavior, β must be negative. An illustration of the softplus function for different values of α and β is shown below in Figure 10.




Figure 10. Softplus barrier function for 3 different parameter pairs

The selection of the parameter values depends on the design of the measurement covariance for obstacle avoidance. The specific parameter values used will be described in Section 4.

Road Boundary Satisfaction

For the driving requirement of staying within the road boundaries, another barrier function is used. The input to this barrier function can be calculated in a similar manner to the lane centerline deviation for the straight road case. The global y-coordinate of the particle state can be compared to the y-coordinate of the road boundaries, and if the ego vehicle's y-coordinate is outside of the road boundary, the measurement function will provide almost zero weight to that particle due to the strict parameter design of this driving requirement. This road boundary satisfaction measurement is an extra safety mechanism, as the lane centerline and heading angle tracking should only generate trajectories that stay within the road boundary. The specific equation for this barrier function is

$$\frac{1}{m} \ln(1 + e^{n(|y|-w)}), \quad (50)$$

where the parameters are m, n and the lane width w appears as an offset to shift the barrier away from the origin. Note that this implementation requires the road centerline to be at $x = 0$. The plotted function is below in Figure 11.




Figure 11. Barrier function (50) illustration for road boundary satisfaction. The solid line is the barrier function. The dashed line is the road boundary.

Heading Angle Tracking

The requirement to keep the ego vehicle's heading angle close to the road angle is to combat overshooting the target lane during a lane change. This problem is illustrated below in Figure 12.




Figure 12. Heading angle tracking requirement illustration

On the left of Figure 12 is a vehicle state that would have a high weight if this driving requirement was not considered. The vehicle is exactly on the lane centerline, but its velocity direction will cause it to overshoot the reference in the next time step, no matter what inputs are chosen. By implementing a heading angle tracking term, the vehicle state on the right will have a higher weight. The heading angle and lane centerline tracking are at odds with each other during a lane change, but if the parameters associated with these terms in the measurement function and measurement noise covariance \mathbf{R} are designed properly, then smooth lane-changing trajectories will be generated with minimal overshoot. The heading angle error is measured by taking direct state feedback of the vehicle's heading angle and evaluating how well it satisfies the requirement of zero deviation from road angle. The vehicle heading angle is illustrated in Figure 13 below.




Figure 13. Vehicle heading angle illustration

C. Driving Mode Design and Selection

For each driving mode, the covariance matrices \mathbf{Q} and \mathbf{R} can be adjusted to increase trajectory quality over that of trajectories generated with static covariances. The adjustments can be explained intuitively by understanding what each driving mode is trying to accomplish.

The lane-keeping driving mode is the default mode that should be selected for most of the time. When lane-keeping in a straight road scenario absent of obstacles, the vehicle does not need to make any turning maneuvers except for correcting for small errors. Thus, the covariance for the steering rate can be decreased in the input noise covariance matrix. Following the same logic, the measurement covariance matrix elements corresponding to the lane centerline and heading angle tracking can also be reduced. This tightens up the lateral range of the generated trajectories and the result will track straighter than it would with “looser” covariance values. For lane-changing modes, the vehicle is trying to turn out of its current lane and track a new reference lane, so these modes require larger steering rate inputs than lane-keeping. For the measurements, the covariance for the lane and heading angle tracking must be increased to allow for sufficient exploration of the input space to reach a region of high enough likelihood for the filtering weights to not collapse.

When sampling the driving mode for the next trajectory, the transition probabilities first need to be calculated. Sampling a driving mode m_j from a transition probability is expressed as

$$m_j \sim p(m_j | \Delta_l) \quad (51)$$

where Δ_l indicates the evaluation of (48) for only the obstacle occupying the same lane as the ego vehicle. This is done to incentivize the lane changing modes to be selected more frequently when the ego vehicle gets closer to the obstacle avoidance region. Intuitively, one can think about this as modeling the probability that a certain mode will produce a good trajectory given the distance to the leading obstacle.

To implement this, the probability of sampling lane-keeping is first calculated as

$$p(LK) = \begin{cases} \frac{-(P_{base} - P_{min})}{v_x^2} (\Delta_l - v_x)^2 + P_{base}, & 0 < \Delta_l < v_x \\ P_{base}, & \Delta_l \geq v_x \end{cases} \quad (52)$$

There are two parameters, P_{base} and P_{min} in this equation. P_{base} sets the max sampling probability for a lane change and is set at 0.9 for the results presented in Section 4. P_{min} is the lower limit for the lane-keeping sampling probability and is set at 0.1. These parameters can be tuned depending on how many

modes are sampled in each planning phase or based on how large the feasible mode set is. In this study, the feasible modes can be restricted to just two once the ego vehicle's current lane is identified, as the mode that would cause the vehicle to leave the road can be eliminated. Plotting this equation reveals the nature of how the lane-keeping sampling probability is determined as a function of Δ_l and v_x . The plot also demonstrates how changing the parameters affects the probability of sampling a lane change.




Figure 14. Plot of (52) for $v_x = 30 \frac{m}{s}$

In this example, the longitudinal velocity is taken to be 30 m/s. As the distance to the obstacle decreases below the distance the ego vehicle will travel in 1 second, the probability of sampling the lane-keeping mode decreases quadratically to P_{min} . At this point, the sampled trajectories likely include a lane-changing mode that is better than lane-keeping. As the feasible set of driving modes has been pared down to two, if lane-keeping is not sampled, then the other feasible mode will be selected. When the ego vehicle is far from any obstacles in its lane, then the algorithm does not expect a lane-changing mode to provide a better trajectory than lane-keeping. This means that the ego vehicle does not prefer any specific lane for steady state operation, opposed to regular driving laws that require vehicles keep right except when passing.

After the set number of driving modes are sampled and corresponding trajectories generated, the best one is chosen for execution. The method of determining which driving mode to commit to is to implement a cumulative cost function that evaluates each generated trajectory. The trajectories are generated according to the reference signal associated with the driving mode sampled. The cost function is

$$C = \sum_{horizon} \gamma \|e_v\|^2 + \epsilon \|e_l\|^2 + \zeta \ln(1 + e^{\nu(\Delta_{obs} - \xi)}) + \lambda \ln(1 + e^{\eta(|y| - w + \rho)}), \quad (53)$$

where there are several weighting parameters that can be tuned to adjust the cost function's preference towards certain behaviors. For example, the lane boundary satisfaction and obstacle avoidance barrier functions are safety-critical, so their parameters should be designed such that any violation ensures that the trajectory will not be selected. Through the filtering and smoothing, the trajectories should already satisfy the constraints of the environment, but the cost function allows for a second pass over all generated trajectories before final application of inputs to the true system.

The cost function works in a similar manner to the measurement function and evaluates how well a trajectory tracks the nominal velocity, reference lane centerline, and satisfaction of obstacle and road boundary constraints. There is some bias implicit to this formulation and this bias explains why there is no heading angle tracking term in the cost function. This bias appears because at the beginning of a lane-changing trajectory, the lane tracking error will be large (as the vehicle is in the wrong lane) but will decrease towards the end of the trajectory as the states moves towards the target lane. Because the cost function is cumulative over the entire planning horizon, this early error can lead to inflated cost even though the trajectory is good for changing lanes. The heading angle also must change during a lane-changing maneuver, so adding a heading angle penalty term in the cost function would further increase the bias against selecting a lane-changing trajectory.

As seen in (53), the cumulative velocity and lane tracking terms are simply a sum of squared errors. The obstacle avoidance and lane boundary terms are again a softplus barrier function. The cost function can have different parameters than the measurement function and has offset terms applied in the softplus function that are not in the measurement function. The cost function term for obstacle avoidance includes an offset ξ which disincentivizes trajectories that remain near the obstacle region, even if they do not violate the obstacle region. This is effectively inflating the size of the obstacle ellipse and combats lane-keeping trajectories that skirt around the edge of an obstacle instead of committing to a lane change. Without this offset, lane-changing can be selected later than desired and the algorithm is generally more indecisive when a lane change is required.

The lane boundary barrier function includes an offset ρ , which can be thought of as narrowing the available road width. This is illustrated below in Figure 14 where the boundary barrier function is plotted with several values of ρ . This offset is implemented to combat the possibility that a lane-keeping trajectory would be selected that goes between an obstacle ellipse and the road boundary, causing the ego vehicle to get “trapped” in an unsafe region where all trajectories would be costly in the next planning phase.




Figure 15. Cost function 4th term – road boundary softplus function with varying offset parameters. The gray vertical lines indicate the road boundaries. Other parameters are $\lambda = 5, \eta = 15$ for all plots.

The overall motion planning algorithm is described below.

Algorithm 3 – PF-RS Motion Planning

Algorithm describes the process for a single planning phase

- Observe road information (obstacle positions and velocities, current ego vehicle lane)
- Sample M_{max} driving modes with (52)
- For each sampled driving mode
 - Identify reference lane, input noise covariance, and measurement noise covariance
 - Execute forward filtering with Algorithm 1
 - Execute reweighting smoothing with Algorithm 2
 - Evaluate the cumulative cost of the generated trajectory with (53)
- End for
- Choose best trajectory (minimum cost)
- Apply 1st set of inputs from best trajectory to the true system

4. Simulation Study

Single Planning Horizon Simulation

In this section, the simulation results will be built up to illustrate what is happening in each step and how certain parameters effect the resulting trajectories. Because there is random sampling involved, the results shown here will all be generated with the MATLAB random number generation seed set to 'default'. This will ensure consistent results when the simulation is run multiple times.

First, the particles themselves are illustrated in increasing quantities. As the particle count increases, it becomes harder to distinguish them, but the exploration of the state space becomes much better and higher quality trajectories are the result. In the lane-keeping driving mode, the vehicle is trying to maintain its lane center tracking for its current lane, and for a straight road this results in a straight trajectory with nominally zero steering input. The vehicle is still trying to maintain the nominal velocity (speed limit) as well. For the lane-changing driving mode, the vehicle tracks the nominal velocity but must make steering corrections to track a different lane center than that of its currently occupied lane.

Figures 15 and 16 on the following pages demonstrate how varying the particle count affects the generated trajectory for lane-keeping and lane-changing, respectively. Individual particles are plotted along with the resulting weighted sum trajectory calculated with (38). Note that these trajectories are after using particle filtering and reweighting particle smoothing. The resampling steps are occurring whenever the particles collapse back into a smaller area. All the lane-keeping simulated trajectories have the same process and measurement covariances. All the lane-changing trajectories have the same process and measurement covariances. The black **x** at the left end of each trajectory is the initial position of the ego vehicle.




Figure 15. Individual particle evolution is plotted in the upper figures, with the resulting trajectory plotted below. (a) 50 particles, (b) 100 particles, and (c) 500 particles.




Figure 16. Individual particle evolution is plotted in the upper figures, with the resulting trajectory plotted below. (a) 50 particles, (b) 100 particles, and (c) 500 particles.

For all driving modes, the state variable process noise covariance \mathbf{Q}_0 is kept constant and is given below in (54). The subscript indicates the matrix dimension.

$$\mathbf{Q}_{0_{n_x \times n_x}} = \begin{bmatrix} 1 \times 10^{-4} & 0 & 0 & 0 & 0 \\ 0 & 1 \times 10^{-4} & 0 & 0 & 0 \\ 0 & 0 & 1 \times 10^{-7} & 0 & 0 \\ 0 & 0 & 0 & 1 \times 10^{-4} & 0 \\ 0 & 0 & 0 & 0 & 1 \times 10^{-7} \end{bmatrix} \quad (54)$$

This matrix is designed to give the system increased confidence in the steering angle and heading angle states, each of which can have detrimental effects on the trajectory tracking if the initial state estimate is slightly deviated from the true state. The input noise covariance matrix is varied for each driving mode, and the specific values used in the simulation study are shown in Table 1. The values reflect the goals of each mode. Compared to lane-keeping, the lane-changing modes have an increased steering rate noise covariance to allow for more steering exploration which results in more lateral movement. The lane tracking and heading angle tracking covariances are increased accordingly in the measurement noise covariance matrix. These changes reduce the weight penalty associated with leaving the ego vehicle's current lane.

Table 1. Input and measurement noise covariance matrices applied in simulations.

| Mode | Input Noise Covariance | Measurement Noise Covariance |
|---------------|--|--|
| Lane-Keeping | $\mathbf{Q}_{n_u \times n_u} = \begin{bmatrix} 1 & 0 \\ 0 & 0.005 \end{bmatrix}$ | $\mathbf{R} = \begin{bmatrix} 2 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0.025 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0.1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0.1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0.0002 \end{bmatrix}$ |
| Lane-Changing | $\mathbf{Q}_{n_u \times n_u} = \begin{bmatrix} 1 & 0 \\ 0 & 0.01 \end{bmatrix}$ | $\mathbf{R} = \begin{bmatrix} 2 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0.1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0.1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0.1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0.0007 \end{bmatrix}$ |

After constructing the virtual system process noise covariance, a multiplier can be applied to increase the covariance for large particle counts. This helps with the estimation quality and improves numerical stability, as the region of high likelihood is larger, and more particles can be used to get a good estimate of the posterior. For 500 particles, this multiplier is set to $G = 10$.

$$\mathbf{Q}' = G \times \mathbf{Q}, \mathbf{R}' = G \times \mathbf{R} \quad (55)$$

The driving scenario analyzed in the full-length simulations has two obstacles; one obstacle occupies each lane. The right lane obstacle is traveling at 15 m/s and the left lane obstacle is traveling at 17 m/s. The ego vehicle starts the simulation at 20 m/s and must speed up to the nominal velocity of 30 m/s. The obstacle positions and their elliptical avoidance regions are shown below. Figures 17 and 18 display the

five generated trajectories in two different cases. The first case in Figure 17 is when the lane-keeping mode sampling probability is 90% due to the distance to the red obstacle being greater than the distance the vehicle will travel in 1 second. The second case in Figure 18 is when the lane-keeping mode sampling probability is reduced due to the ego vehicle being closer to the obstacle.




Figure 17. Scenario layout with 5 generated trajectories (4 LK, 1 CLL) with LK sampling probability of 90%.




Figure 18. Scenario with ego vehicle adjusted to show impact of (52). of the 5 generated trajectories, 3 are LK and 2 are CLL. The LK sampling probability is 64.6%.




Figure 19. Zoomed in view of generated trajectories in Figure 18.

In Figures 18 and 19 above, it may seem that the lane-keeping trajectories are violating the obstacle ellipse, however this is not the case. Recall that during the planning phase, the obstacle position is predicted based on its measured velocity. As the obstacle moves, the elliptical region will move with it, and the trajectories shown do not violate the region at the time when they coexist. Next, the results of full-length simulations are presented.

The first simulation was run with 250 particles and 5 driving modes sampled per planning phase ($M_{max} = 5$). The plots below show the lane tracking and velocity tracking. In Figure 20, the target lane centerline is shown as the dotted line. There are some instances of unexpected or inconsistent mode choices that occur. These happen during the transient phases of the ego vehicles motion when deciding to start a lane change. Sometimes, just through random chance and the MATLAB random number generation method, only 1 driving mode will be sampled 5 times. This can cause lane-keeping to be selected during a lane-changing maneuver, which appears in these plots. Once the ego vehicle gets close

enough to the leading obstacle for the lane-keeping sampling probability to reach P_{min} , the chances of sampling all 5 modes as lane-keeping diminishes to 0.001% as per the binomial distribution with $n = 5, p = 0.1, x = 5$. Figure 20 shows the lane tracking performance versus longitudinal position and versus time.




Figure 20. Lane centerlines of driving mode and vehicle position. 250 particles, 5 modes sampled, lane-keeping sampling probability parameters $P_{base} = 0.9, P_{min} = 0.1$

The inputs and the corresponding state variables over the runtime of the simulation are shown below in Figures 21 and 22.




Figure 21. Acceleration inputs and corresponding velocity vs time.




Figure 22. Steering rate inputs and corresponding angles vs time.

In these 250-particle simulations, the estimated trajectories are somewhat jagged, although the driving requirements are still satisfied. The next simulation has the same setup, except 500 particles are used. In these simulations, it is expected that increased particle count will improve the estimation performance. Again, the results for lane tracking are shown first with the position and time axis in Figure 23.




Figure 23. Lane centerlines of driving mode and vehicle position. 500 particles, 5 modes sampled, lane-keeping sampling probability parameters $P_{base} = 0.9, P_{min} = 0.1$

The inputs and the corresponding state variables are shown below in Figures 24 and 25.




Figure 24. Acceleration inputs and corresponding velocity vs time.




Figure 25. Steering rate inputs and corresponding angles vs time.

When these two sets of results are compared, the improvement in performance by doubling the particle count can be seen. The vehicle's behavior is overall smoother, as seen in the lane tracking graph and the smaller changes in velocity (smaller accelerations). The nominal velocity tracking is improved after finishing the initial acceleration. The steering behavior also becomes more aligned with what a human driver would do during a lane change, with a clearly visible turn left - turn right sequence between 2 and 4 seconds corresponding to a lane change to the left. The opposite is clearly visible between 7 and 9 seconds for a lane change to the right. Animated simulation results are available in the GitHub repository for this project [8].

5. Conclusion

This study has presented a particle filter and reweighting particle smoother-based motion planning method for autonomous road vehicles in a dynamic highway driving environment. The generated trajectories are dynamically feasible though they may be suboptimal. Barrier functions in the particle filter weight update equations and trajectory selection cost function are used to prevent collisions or road boundary violation. The inputs extracted from the generated trajectories are applied to the system in a receding horizon manner, which enables responsiveness to the dynamic environment.

Applying this method to more sophisticated scenarios requires further work, as the formulation of the obstacle regions and driving mode selection will need reworking for curved roads or more than two lanes. A potential improvement to the driving modes would be an implementation of a “following” mode for cases when there is no open lane to make a passing maneuver. This mode would set the nominal velocity of the ego vehicle to the velocity of the leading vehicle and track the current lane centerline. The driving mode sampling could also be modified to be conditioned on the current mode as in [7].

Potential improvements to the trajectory generation method of forward particle filtering and reweighting particle smoothing can come from choosing better proposal distributions to improve the quality of particles sampled at each time step. There are variations of the particle filter that seek to generate more optimal proposal distributions for nonlinear systems than the bootstrap particle filter, which may be worth exploration. This method could potentially be expanded to motion planning for multiple vehicles in a collaborative manner to enable platooning or cooperative merging.

6. References

- [1] C. Kuptamettee and N. Aunsri, "A review of resampling techniques in particle filtering framework," *Measurement*, vol. 193, 2022, 110836, doi: 10.1016/j.measurement.2022.110836.
- [2] D. González, J. Pérez, V. Milanés and F. Nashashibi, "A Review of Motion Planning Techniques for Automated Vehicles," in *IEEE Transactions on Intelligent Transportation Systems*, vol. 17, no. 4, pp. 1135-1145, April 2016, doi: 10.1109/TITS.2015.2498841.
- [3] D. Stahl and J. Hauth, "PF-MPC: Particle filter-model predictive control," *Systems & Control Letters*, vol. 60, no. 8, pp. 632-643, 2011
- [4] F. Farshidian, E. Jelavic, A. Satapathy, M. Gifftthaler and J. Buchli, "Real-time motion planning of legged robots: A model predictive control approach," *2017 IEEE-RAS 17th International Conference on Humanoid Robotics (Humanoids)*, 2017, pp. 577-584, doi: 10.1109/HUMANOIDS.2017.8246930.
- [5] I. Askari, S. Zeng and H. Fang, "Nonlinear Model Predictive Control Based on Constraint-Aware Particle Filtering/Smoothing," *2021 American Control Conference (ACC)*, 2021, pp. 3532-3537, doi: 10.23919/ACC50511.2021.9482774.
- [6] K. Berntorp, P. Inani, R. Quirynen and S. Di Cairano, "Motion Planning of Autonomous Road Vehicles by Particle Filtering: Implementation and Validation," *2019 American Control Conference (ACC)*, 2019, pp. 1382-1387, doi: 10.23919/ACC.2019.8815309.
- [7] K. Berntorp, T. Hoang and S. Di Cairano, "Motion Planning of Autonomous Road Vehicles by Particle Filtering," in *IEEE Transactions on Intelligent Vehicles*, vol. 4, no. 2, pp. 197-210, June 2019, doi: 10.1109/TIV.2019.2904394.
- [8] M. Lieb, PF-RS Motion Planning, 2022, GitHub repository, <https://github.com/MasonL-KU/PF-RS-Motion-Planning>
- [9] R. van der Merwe, A. Doucet, J. F. G. de Freitas, E. Wan, "The Unscented Particle Filter", Technical Report CUED/F-INFENG/TR 380, 2000, Cambridge University Engineering Department.
- [10] S. Karaman, M. R. Walter, A. Perez, E. Frazzoli and S. Teller, "Anytime Motion Planning using the RRT*," *2011 IEEE International Conference on Robotics and Automation*, 2011, pp. 1478-1483, doi: 10.1109/ICRA.2011.5980479.
- [11] S. M. LaValle and J. J. Kuffner, "Randomized kinodynamic planning," *Proceedings 1999 IEEE International Conference on Robotics and Automation (Cat. No.99CH36288C)*, 1999, pp. 473-479 vol.1, doi: 10.1109/ROBOT.1999.770022.
- [12] S. Särkkä (2013). *Bayesian Filtering and Smoothing* (Institute of Mathematical Statistics Textbooks). Cambridge: Cambridge University Press. doi:10.1017/CBO9781139344203

- [13] X. Qian, I. Navarro, A. de La Fortelle and F. Moutarde, "Motion planning for urban autonomous driving using Bézier curves and MPC," *2016 IEEE 19th International Conference on Intelligent Transportation Systems (ITSC)*, 2016, pp. 826-833, doi: 10.1109/ITSC.2016.7795651.
- [14] Y. Kuwata, G. A. Fiore, J. Teo, E. Frazzoli and J. P. How, "Motion planning for urban driving using RRT," *2008 IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2008, pp. 1681-1686, doi: 10.1109/IROS.2008.4651075.

Appendix

Equation (22) Derivation

Markov properties of the model for developing the sequential importance sampling equations.

$$\mathbf{x}_k \sim p(\mathbf{x}_k | \mathbf{x}_{k-1})$$

$$\mathbf{y}_k \sim p(\mathbf{y}_k | \mathbf{x}_k)$$

Create recursive estimate for the full posterior $p(\mathbf{x}_{0:k} | \mathbf{y}_{1:k})$. By applying the definition of conditional probability, this can be rewritten as follows.

$$\begin{aligned} p(\mathbf{x}_{0:k} | \mathbf{y}_{1:k}) &= \frac{p(\mathbf{y}_k, \mathbf{y}_{1:k-1}, \mathbf{x}_{0:k})}{p(\mathbf{y}_k, \mathbf{y}_{1:k-1})} = \frac{p(\mathbf{y}_k | \mathbf{x}_{0:k}, \mathbf{y}_{1:k-1}) p(\mathbf{x}_{0:k} | \mathbf{y}_{1:k-1}) p(\mathbf{y}_{1:k-1})}{p(\mathbf{y}_k | \mathbf{y}_{1:k-1}) p(\mathbf{y}_{1:k-1})} \\ p(\mathbf{x}_{0:k} | \mathbf{y}_{1:k}) &= \frac{p(\mathbf{y}_k | \mathbf{x}_{0:k}, \mathbf{y}_{1:k-1}) p(\mathbf{x}_{0:k} | \mathbf{y}_{1:k-1})}{p(\mathbf{y}_k | \mathbf{y}_{1:k-1})} \end{aligned}$$

The denominator can be removed by taking a proportionality.

$$p(\mathbf{x}_{0:k} | \mathbf{y}_{1:k}) \propto p(\mathbf{y}_k | \mathbf{x}_{0:k}, \mathbf{y}_{1:k-1}) p(\mathbf{x}_{0:k} | \mathbf{y}_{1:k-1})$$

Expanding the first term in this equation and applying Markov properties of the system.

$$\begin{aligned} p(\mathbf{y}_k | \mathbf{x}_{0:k}, \mathbf{y}_{1:k-1}) &= \frac{p(\mathbf{y}_k, \mathbf{x}_{0:k}, \mathbf{y}_{1:k-1})}{p(\mathbf{x}_{0:k}, \mathbf{y}_{1:k-1})} \\ &= \frac{p(\mathbf{y}_k | \mathbf{x}_{0:k}, \mathbf{y}_{1:k-1}) p(\mathbf{x}_k | \mathbf{x}_{0:k-1}, \mathbf{y}_{1:k-1}) p(\mathbf{x}_{0:k-1} | \mathbf{y}_{1:k-1}) p(\mathbf{y}_{1:k-1})}{p(\mathbf{x}_{0:k} | \mathbf{y}_{1:k-1}) p(\mathbf{y}_{1:k-1})} \\ p(\mathbf{y}_k | \mathbf{x}_{0:k}, \mathbf{y}_{1:k-1}) &= \frac{p(\mathbf{y}_k | \mathbf{x}_{0:k}, \mathbf{y}_{1:k-1}) p(\mathbf{x}_k | \mathbf{x}_{0:k-1}, \mathbf{y}_{1:k-1}) p(\mathbf{x}_{0:k-1} | \mathbf{y}_{1:k-1})}{p(\mathbf{x}_{0:k} | \mathbf{y}_{1:k-1})} \end{aligned}$$

Plugging this back into the equation and applying the Markov property for the states and measurements.

$$\begin{aligned} p(\mathbf{y}_k | \mathbf{x}_{0:k}, \mathbf{y}_{1:k-1}) &= p(\mathbf{y}_k | \mathbf{x}_k) p(\mathbf{x}_k | \mathbf{x}_{0:k-1}, \mathbf{y}_{1:k-1}) p(\mathbf{x}_{0:k-1} | \mathbf{y}_{1:k-1}) \\ &= p(\mathbf{y}_k | \mathbf{x}_k) p(\mathbf{x}_k | \mathbf{x}_{k-1}) p(\mathbf{x}_{0:k-1} | \mathbf{y}_{1:k-1}) \end{aligned}$$

Now, the result of (22) has been reached.

Equations (23)-(25) Derivation

These equations develop the recursive weight update for the sequential importance sampling algorithm. First, the recursive property of the proposal distribution $\pi(\mathbf{x}_{0:k}^i | \mathbf{y}_{1:k})$ is demonstrated.

$$\pi(\mathbf{x}_{0:k}^i | \mathbf{y}_{1:k}) = \frac{\pi(\mathbf{x}_{0:k}^i, \mathbf{y}_{1:k})}{\pi(\mathbf{y}_{1:k})} = \frac{\pi(\mathbf{x}_k^i | \mathbf{x}_{0:k-1}^i, \mathbf{y}_{1:k}) \pi(\mathbf{x}_{0:k-1}^i | \mathbf{y}_{1:k}) \pi(\mathbf{y}_{1:k})}{\pi(\mathbf{y}_{1:k})}$$

$$\pi(\mathbf{x}_{0:k}^i | \mathbf{y}_{1:k}) = \pi(\mathbf{x}_k^i | \mathbf{x}_{0:k-1}^i, \mathbf{y}_{1:k}) \times \pi(\mathbf{x}_{0:k-1}^i | \mathbf{y}_{1:k})$$

We enforce the following property on the proposal distribution to facilitate the recursive weight update.

$$\pi(\mathbf{x}_{0:k-1}^i | \mathbf{y}_{1:k}) = \pi(\mathbf{x}_{0:k-1}^i | \mathbf{y}_{1:k-1})$$

This property essentially says that particles at time k are independent of the future measurement \mathbf{y}_{k+1} . Applying this proposal distribution to the weight update and including the expanded proposal, arrive at (23)

$$w_k^i \propto \frac{p(\mathbf{y}_k | \mathbf{x}_k^i) p(\mathbf{x}_k^i | \mathbf{x}_{k-1}^i) p(\mathbf{x}_{0:k-1}^i | \mathbf{y}_{1:k-1})}{\pi(\mathbf{x}_{0:k}^i | \mathbf{y}_{1:k})} = \frac{p(\mathbf{y}_k | \mathbf{x}_k^i) p(\mathbf{x}_k^i | \mathbf{x}_{k-1}^i)}{\pi(\mathbf{x}_k^i | \mathbf{x}_{0:k-1}^i, \mathbf{y}_{1:k})} \times \frac{p(\mathbf{x}_{0:k-1}^i | \mathbf{y}_{1:k-1})}{\pi(\mathbf{x}_{0:k-1}^i | \mathbf{y}_{1:k-1})}$$

Observing that the weight for the previous estimate of the posterior at time $k - 1$ is presented inside of (23), arrive at (25).

$$w_k^i \propto \frac{p(\mathbf{y}_k | \mathbf{x}_k^i) p(\mathbf{x}_k^i | \mathbf{x}_{k-1}^i)}{\pi(\mathbf{x}_k^i | \mathbf{x}_{0:k-1}^i, \mathbf{y}_{1:k})} \times w_{k-1}^i$$

Equation (33) Derivation

The derivation starts from the Bayesian fixed interval smoothing equations.

$$p(\mathbf{x}_{k+1}|\mathbf{y}_{1:k}) = \int p(\mathbf{x}_{k+1}|\mathbf{x}_k)p(\mathbf{x}_k|\mathbf{y}_{1:k})d\mathbf{x}_k$$

$$p(\mathbf{x}_k|\mathbf{y}_{1:T}) = p(\mathbf{x}_k|\mathbf{y}_{1:k}) \int \left[\frac{p(\mathbf{x}_{k+1}|\mathbf{x}_k)p(\mathbf{x}_{k+1}|\mathbf{y}_{1:T})}{p(\mathbf{x}_{k+1}|\mathbf{y}_{1:k})} \right] d\mathbf{x}_{k+1}$$

From the forward filtering stage, the marginal filtering distributions are already calculated.

$$p(\mathbf{x}_{k+1}|\mathbf{y}_{1:T}) \approx \sum_i w_{k+1|T}^i \delta(\mathbf{x}_{k+1} - \mathbf{x}_{k+1}^i)$$

With this approximation available, the integral term can be approximated.

$$\int \left[\frac{p(\mathbf{x}_{k+1}|\mathbf{x}_k)p(\mathbf{x}_{k+1}|\mathbf{y}_{1:T})}{p(\mathbf{x}_{k+1}|\mathbf{y}_{1:k})} \right] d\mathbf{x}_{k+1} \approx \int \left[\frac{p(\mathbf{x}_{k+1}|\mathbf{x}_k)}{p(\mathbf{x}_{k+1}|\mathbf{y}_{1:k})} \right] \sum_i [w_{k+1|T}^i \delta(\mathbf{x}_{k+1} - \mathbf{x}_{k+1}^i)] d\mathbf{x}_{k+1}$$

$$= \sum_i w_{k+1|T}^i \frac{p(\mathbf{x}_{k+1}^i|\mathbf{x}_k)}{p(\mathbf{x}_{k+1}^i|\mathbf{y}_{1:k})}$$

This is by applying the following property of the Dirac delta function.

$$\int_{-\infty}^{\infty} f(x)\delta(x-a)dx = f(a)$$

The denominator term $p(\mathbf{x}_{k+1}|\mathbf{y}_{1:k})$ is the predictive distribution also approximated using the filtering method.

$$p(\mathbf{x}_{k+1}|\mathbf{y}_{1:k}) \approx \sum_j w_k^j p(\mathbf{x}_{k+1}|\mathbf{x}_k^j)$$

With this approximation, arrive at the following approximation for the integral.

$$\int \left[\frac{p(\mathbf{x}_{k+1}|\mathbf{x}_k)p(\mathbf{x}_{k+1}|\mathbf{y}_{1:T})}{p(\mathbf{x}_{k+1}|\mathbf{y}_{1:k})} \right] d\mathbf{x}_{k+1} \approx \sum_i w_{k+1|T}^i \frac{p(\mathbf{x}_{k+1}^i|\mathbf{x}_k)}{[\sum_j w_k^j p(\mathbf{x}_{k+1}^i|\mathbf{x}_k^j)]}$$

Returning to the overall Bayesian optimal smoothing equation and substituting this approximation, arrive at the following.

$$p(\mathbf{x}_k|\mathbf{y}_{1:k}) \int \left[\frac{p(\mathbf{x}_{k+1}|\mathbf{x}_k)p(\mathbf{x}_{k+1}|\mathbf{y}_{1:T})}{p(\mathbf{x}_{k+1}|\mathbf{y}_{1:k})} \right] d\mathbf{x}_{k+1} \approx p(\mathbf{x}_k|\mathbf{y}_{1:k}) \sum_i w_{k+1|T}^i \frac{p(\mathbf{x}_{k+1}^i|\mathbf{x}_k)}{[\sum_j w_k^j p(\mathbf{x}_{k+1}^i|\mathbf{x}_k^j)]}$$

Where $p(\mathbf{x}_k|\mathbf{y}_{1:k})$ is again approximated as

$$p(\mathbf{x}_k|\mathbf{y}_{1:k}) \approx \sum_l w_k^l \delta(\mathbf{x}_k - \mathbf{x}_k^l)$$

Substituting this approximation, arrive at the following.

$$p(\mathbf{x}_k | \mathbf{y}_{1:k}) \sum_i w_{k+1|T}^i \frac{p(\mathbf{x}_{k+1}^i | \mathbf{x}_k)}{[\sum_j w_k^j p(\mathbf{x}_{k+1}^j | \mathbf{x}_k^j)]} \approx \sum_l w_k^l \delta(\mathbf{x}_k - \mathbf{x}_k^l) \sum_i w_{k+1|T}^i \frac{p(\mathbf{x}_{k+1}^i | \mathbf{x}_k^l)}{[\sum_j w_k^j p(\mathbf{x}_{k+1}^j | \mathbf{x}_k^j)]}$$

Finally, the weight update equation for the smoothing weights for particle i at time k is reached.

$$p(\mathbf{x}_k | \mathbf{y}_{1:T}) \approx \sum_l w_k^l \delta(\mathbf{x}_k - \mathbf{x}_k^l) \sum_i w_{k+1|T}^i \frac{p(\mathbf{x}_{k+1}^i | \mathbf{x}_k^l)}{[\sum_j w_k^j p(\mathbf{x}_{k+1}^j | \mathbf{x}_k^j)]}$$

Combining all the terms except the $\delta(\cdot)$ gives the simplified approximation to the smoothing distribution

$$p(\mathbf{x}_k | \mathbf{y}_{1:T}) \approx \sum_l w_{k|T}^l \delta(\mathbf{x}_k - \mathbf{x}_k^l)$$

where the weights are updated backwards using

$$w_{k|T}^l = \sum_i w_{k+1|T}^i \frac{w_k^l p(\mathbf{x}_{k+1}^i | \mathbf{x}_k^l)}{[\sum_j w_k^j p(\mathbf{x}_{k+1}^j | \mathbf{x}_k^j)]}$$